




Figure 35: Score@k with 30-minute runs for Fix Embedding.




Figure 36: Score@k with 2-hour runs for Fix Embedding.




Figure 37: Score@k with 30-minute runs for Scaling Law Experiment.




Figure 38: Score@k with 2-hour runs for Scaling Law Experiment.




Figure 39: Score@k with 30-minute runs for Restricted Architecture MLM.




Figure 40: Score@k with 2-hour runs for Restricted Architecture MLM.




Figure 41: Score@k with 30-minute runs for Finetune GPT-2 for QA.




Figure 42: Score@k with 2-hour runs for Finetune GPT-2 for QA.




Figure 43: Score@k with 30-minute runs for Scaffolding for Rust CodeContests.




Figure 44: Score@k with 2-hour runs for Scaffolding for Rust CodeContests.

I Score over time by task (8 hours)

Here we present score over time for 8-hour runs for each task. Average performance is represented with darker lines, while the lighter lines represent individual runs.




Figure 45: 8-hour performance for Optimize LLM Foundry.




Figure 46: 8-hour performance for *Optimize a Kernel*.




Figure 47: 8-hour performance for Fix Embedding.